

# Administration

- ▶ Assignment #2 Posted
  - ▶ Due January 17 (Mon.)
- ▶ Midterm exam
  - ▶ January 11 (Tues.)
  - ▶ Open book, Open notes
  - ▶ Will need “calculator”
    - ▶ May use laptop, iPhone, etc.

# Parallel Computing I

SMP: Sequential Dependencies, Barrier Actions, and Overlapping

# Looking Back, Looking Forward

Last three weeks:

- ▶ Why parallel computing?
- ▶ Parallel program designs
- ▶ Massively parallel problems
- ▶ SMP parallel programs with Parallel Java
  - ▶ parallel teams
  - ▶ parallel for loops
- ▶ Performance metrics
- ▶ Load balancing and reduction

This week:

- ▶ sequential dependencies
- ▶ barrier actions
- ▶ overlapping

## Quick Review


- ▶ FindKey
- ▶ MandelbrotSet
- ▶ Pi
- ▶ MSHistogram

### Performance on cadmium

- ▶ Four AMD Opteron 6172 12-core CPUs (48 processors),  
2.1 GHz clock, 128 GB main memory

# All-Pairs Shortest-Path Problem

Given a graph with weighted edges,  
determine the (length of the) shortest path between all pairs of vertices.



	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>	<i>F</i>	<i>G</i>	<i>H</i>	<i>I</i>	<i>J</i>
<i>A</i>	0	$\infty$	$\infty$	462	$\infty$	451	$\infty$	370	$\infty$	$\infty$
<i>B</i>	$\infty$	0	190	$\infty$	$\infty$	399	$\infty$	$\infty$	$\infty$	$\infty$
<i>C</i>	$\infty$	190	0	$\infty$	$\infty$	234	333	366	414	$\infty$
<i>D</i>	462	$\infty$	$\infty$	0	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
<i>E</i>	$\infty$	$\infty$	$\infty$	$\infty$	0	359	394	269	$\infty$	325
<i>F</i>	451	399	234	$\infty$	359	0	239	144	$\infty$	$\infty$
<i>G</i>	$\infty$	$\infty$	333	$\infty$	394	239	0	337	389	$\infty$
<i>H</i>	370	$\infty$	366	$\infty$	269	144	337	0	$\infty$	$\infty$
<i>I</i>	$\infty$	$\infty$	414	$\infty$	$\infty$	$\infty$	389	$\infty$	0	$\infty$
<i>J</i>	$\infty$	$\infty$	$\infty$	$\infty$	325	$\infty$	$\infty$	$\infty$	$\infty$	0

# Floyd's All-Pairs Shortest-Path Algorithm

```
for  $i = 0$  to  $n - 1$ 
  for  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
      // Update the distance from  $r$  to  $c$  via  $i$ .
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

# Floyd's All-Pairs Shortest-Path Algorithm

```
for  $i = 0$  to  $n - 1$ 
  for  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
      // Update the distance from  $r$  to  $c$  via  $i$ .
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

	A	B	C	D	E	F	G	H	I	J
A	0	$\infty$	$\infty$	462	$\infty$	451	$\infty$	370	$\infty$	$\infty$
B	$\infty$	0	190	$\infty$	$\infty$	399	$\infty$	$\infty$	$\infty$	$\infty$
C	$\infty$	190	0	$\infty$	$\infty$	234	333	366	414	$\infty$
D	462	$\infty$	$\infty$	0	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
E	$\infty$	$\infty$	$\infty$	$\infty$	0	359	394	269	$\infty$	325
F	451	399	234	$\infty$	359	0	239	144	$\infty$	$\infty$
G	$\infty$	$\infty$	333	$\infty$	394	239	0	337	389	$\infty$
H	370	$\infty$	366	$\infty$	269	144	337	0	$\infty$	$\infty$
I	$\infty$	$\infty$	414	$\infty$	$\infty$	$\infty$	389	$\infty$	0	$\infty$
J	$\infty$	$\infty$	$\infty$	$\infty$	325	$\infty$	$\infty$	$\infty$	$\infty$	0

Which entry is the first to be changed?

# Floyd's All-Pairs Shortest-Path Algorithm

```
for  $i = 0$  to  $n - 1$ 
  for  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
      // Update the distance from  $r$  to  $c$  via  $i$ .
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

	A	B	C	D	E	F	G	H	I	J
A	0	$\infty$	$\infty$	462	$\infty$	451	$\infty$	370	$\infty$	$\infty$
B	$\infty$	0	190	$\infty$	$\infty$	399	$\infty$	$\infty$	$\infty$	$\infty$
C	$\infty$	190	0	$\infty$	$\infty$	234	333	366	414	$\infty$
D	462	$\infty$	$\infty$	0	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
E	$\infty$	$\infty$	$\infty$	$\infty$	0	359	394	269	$\infty$	325
F	451	399	234	$\infty$	359	0	239	144	$\infty$	$\infty$
G	$\infty$	$\infty$	333	$\infty$	394	239	0	337	389	$\infty$
H	370	$\infty$	366	$\infty$	269	144	337	0	$\infty$	$\infty$
I	$\infty$	$\infty$	414	$\infty$	$\infty$	$\infty$	389	$\infty$	0	$\infty$
J	$\infty$	$\infty$	$\infty$	$\infty$	325	$\infty$	$\infty$	$\infty$	$\infty$	0

Which entry is the first to be changed?



# Floyd's All-Pairs Shortest-Path Algorithm

```
for  $i = 0$  to  $n - 1$ 
  for  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
      // Update the distance from  $r$  to  $c$  via  $i$ .
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

	A	B	C	D	E	F	G	H	I	J
A	0	$\infty$	$\infty$	462	$\infty$	451	$\infty$	370	$\infty$	$\infty$
B	$\infty$	0	190	$\infty$	$\infty$	399	$\infty$	$\infty$	$\infty$	$\infty$
C	$\infty$	190	0	$\infty$	$\infty$	234	333	366	414	$\infty$
D	462	$\infty$	$\infty$	0	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
E	$\infty$	$\infty$	$\infty$	$\infty$	0	359	394	269	$\infty$	325
F	451	399	234	$\infty$	359	0	239	144	$\infty$	$\infty$
G	$\infty$	$\infty$	333	$\infty$	394	239	0	337	389	$\infty$
H	370	$\infty$	366	$\infty$	269	144	337	0	$\infty$	$\infty$
I	$\infty$	$\infty$	414	$\infty$	$\infty$	$\infty$	389	$\infty$	0	$\infty$
J	$\infty$	$\infty$	$\infty$	$\infty$	325	$\infty$	$\infty$	$\infty$	$\infty$	0

Which entry is the first to be changed?

# Floyd's All-Pairs Shortest-Path Algorithm

```
for  $i = 0$  to  $n - 1$ 
  for  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
      // Update the distance from  $r$  to  $c$  via  $i$ .
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

	A	B	C	D	E	F	G	H	I	J
A	0	$\infty$	$\infty$	462	$\infty$	451	$\infty$	370	$\infty$	$\infty$
B	$\infty$	0	190	$\infty$	$\infty$	399	$\infty$	$\infty$	$\infty$	$\infty$
C	$\infty$	190	0	$\infty$	$\infty$	234	333	366	414	$\infty$
D	462	$\infty$	$\infty$	0	$\infty$	913	$\infty$	$\infty$	$\infty$	$\infty$
E	$\infty$	$\infty$	$\infty$	$\infty$	0	359	394	269	$\infty$	325
F	451	399	234	$\infty$	359	0	239	144	$\infty$	$\infty$
G	$\infty$	$\infty$	333	$\infty$	394	239	0	337	389	$\infty$
H	370	$\infty$	366	$\infty$	269	144	337	0	$\infty$	$\infty$
I	$\infty$	$\infty$	414	$\infty$	$\infty$	$\infty$	389	$\infty$	0	$\infty$
J	$\infty$	$\infty$	$\infty$	$\infty$	325	$\infty$	$\infty$	$\infty$	$\infty$	0

Which entry is the first to be changed?

# Floyd's All-Pairs Shortest-Path Algorithm

```
for  $i = 0$  to  $n - 1$ 
  for  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
      // Update the distance from  $r$  to  $c$  via  $i$ .
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>	<i>F</i>	<i>G</i>	<i>H</i>	<i>I</i>	<i>J</i>
<i>A</i>	0	850	685	462	639	451	690	370	1079	964
<i>B</i>	850	0	190	1312	758	399	523	543	604	1083
<i>C</i>	685	190	0	1147	593	234	333	366	414	918
<i>D</i>	462	1312	1147	0	1101	913	1152	832	1541	1426
<i>E</i>	639	758	593	1101	0	359	394	269	783	325
<i>F</i>	451	399	234	913	359	0	239	144	628	684
<i>G</i>	690	523	333	1152	394	239	0	337	389	719
<i>H</i>	370	543	366	832	269	144	337	0	726	594
<i>I</i>	1079	604	414	1541	783	628	389	726	0	1108
<i>J</i>	964	1083	918	1426	325	684	719	594	1108	0

In general, original entries may change  
and distance matrix may not be symmetric.

# Input/Output Files

Represent distance matrix as an instance of `edu.rit.io.DoubleMatrixFile`:

## ► Input

```
DoubleMatrixFile dmf = new DoubleMatrixFile();  
DoubleMatrixFile.Reader reader = dmf.prepareToRead (instream);  
reader.read();  
reader.close();  
int R = dmf.getRowCount();  
int C = dmf.getColCount();  
double[][] matrix = dmf.getMatrix();
```

## ► Output

```
double[][] matrix = new double [R] [C];  
DoubleMatrixFile dmf = new DoubleMatrixFile( R, C, matrix);  
DoubleMatrixFile.Writer writer = dmf.prepareToWrite (outstream);  
writer.write();  
writer.close();
```

# FloydRandom.java and FloydSeq.java

```
code/FloydRandom.java  
code/FloydSeq.java
```

## FloydRandom.java and FloydSeq.java

code/FloydRandom.java  
code/FloydSeq.java

```
for (int i = 0; i < n; ++i) { double[] d_i = d[i];  
    for (int r = 0; r < n; ++r) { double[] d_r = d[r];  
        for (int c = 0; c < n; ++c) {  
            d_r[c] = Math.min (d_r[c], d_r[i] + d_i[c]);  } } }
```

VS.

```
for (int i = 0; i < n; ++i) {  
    for (int r = 0; r < n; ++r) {  
        for (int c = 0; c < n; ++c) {  
            d[r][c] = Math.min (d[r][c], d[r][i] + d[i][c]);  } } }
```

# Parallelizing Floyd's Algorithm

How do we convert the sequential Floyd program to a parallel Floyd program?

- ▶ What portions of the Floyd program are not parallelizable?
- ▶ What portions of the Floyd program might be parallelizable?

# Parallelizing Floyd's Algorithm

Which loops are parallelizable?

$$d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$$

- ▶ **for** (**int**  $i = 0$ ;  $i < n$ ;  $++i$ ) { ... }
- ▶ On each iteration, store a value into every  $d_{rc}$  that depends upon the values of  $d_{rc}$ ,  $d_{ri}$ , and  $d_{ic}$ , any of which could have been changed on the *previous* iteration.
- ▶ There is a *sequential dependency* from each iteration  $i$  to the next.
- ▶ This loop cannot be parallelized.



# Parallelizing Floyd's Algorithm

Which loops are parallelizable?

$$d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$$

- ▶ **for** (**int**  $r = 0$ ;  $r < n$ ;  $++r$ ) { ... }
- ▶ On each iteration, store a value into every  $d_{rc}$  that depends upon the values of  $d_{rc}$ ,  $d_{ri}$ , and  $d_{ic}$ .
- ▶ Updates to  $d_{ic}$  (when  $r = i$ ) will affect updates to  $d_{rc}$  (when  $r > i$ ).

# Parallelizing Floyd's Algorithm

Which loops are parallelizable?

$$d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$$

- ▶ **for** (**int**  $r = 0$ ;  $r < n$ ;  $++r$ ) { ... }
- ▶ On each iteration, store a value into every  $d_{rc}$  that depends upon the values of  $d_{rc}$ ,  $d_{ri}$ , and  $d_{ic}$ .
- ▶ Updates to  $d_{ic}$  (when  $r = i$ ) will affect updates to  $d_{rc}$  (when  $r > i$ ).
- ▶ But, when  $r = i$ :  $d_{ri} \leftarrow \min(d_{ic}, d_{ii} + d_{ic})$
- ▶ Assuming  $d_{ii} = 0$ , the updates to  $d_{ic}$  (when  $r = i$ ) are idempotent.
- ▶ This loop can be parallelized.

# Parallelizing Floyd's Algorithm

Which loops are parallelizable?

$$d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$$

- ▶ **for** (**int**  $r = 0$ ;  $r < n$ ;  $++r$ ) { ... }
- ▶ On each iteration, store a value into every  $d_{rc}$  that depends upon the values of  $d_{rc}$ ,  $d_{ri}$ , and  $d_{ic}$ .
- ▶ Updates to  $d_{ic}$  (when  $r = i$ ) will affect updates to  $d_{rc}$  (when  $r > i$ ).
- ▶ But, when  $r = i$ :  $d_{ri} \leftarrow \min(d_{ic}, d_{ii} + d_{ic})$
- ▶ Assuming  $d_{ii} = 0$ , the updates to  $d_{ic}$  (when  $r = i$ ) are idempotent.
- ▶ This loop can be parallelized.
- ▶ Any synchronization issues?

# Parallelizing Floyd's Algorithm

Which loops are parallelizable?

$$d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$$

- ▶ **for** (**int**  $r = 0$ ;  $r < n$ ;  $++r$ ) { ... }
- ▶ On each iteration, store a value into every  $d_{rc}$  that depends upon the values of  $d_{rc}$ ,  $d_{ri}$ , and  $d_{ic}$ .
- ▶ Updates to  $d_{ic}$  (when  $r = i$ ) will affect updates to  $d_{rc}$  (when  $r > i$ ).
- ▶ But, when  $r = i$ :  $d_{ri} \leftarrow \min(d_{ic}, d_{ii} + d_{ic})$
- ▶ Assuming  $d_{ii} = 0$ , the updates to  $d_{ic}$  (when  $r = i$ ) are idempotent.
- ▶ This loop can be parallelized.
- ▶ Any synchronization issues?
  - ▶ Java does not guarantee atomicity of reads and writes of a **double**.

# Parallelizing Floyd's Algorithm

Which loops are parallelizable?

$$d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$$


- ▶ **for** (**int**  $r = 0$ ;  $r < n$ ;  $++r$ ) { ... }
- ▶ On each iteration, store a value into every  $d_{rc}$  that depends upon the values of  $d_{rc}$ ,  $d_{ri}$ , and  $d_{ic}$ .
- ▶ Updates to  $d_{ic}$  (when  $r = i$ ) will affect updates to  $d_{rc}$  (when  $r > i$ ).
- ▶ But, when  $r = i$ :  $d_{ri} \leftarrow \min(d_{ic}, d_{ii} + d_{ic})$
- ▶ Assuming  $d_{ii} = 0$ , the updates to  $d_{ic}$  (when  $r = i$ ) are idempotent.
- ▶ This loop can be parallelized.
- ▶ Any synchronization issues?
  - ▶ Java does not guarantee atomicity of reads and writes of a **double**.
  - ▶ "For the purposes of the Java programming language memory model, a single write to a non-volatile long or double value is treated as two separate writes: one to each 32-bit half. This can result in a situation where a thread sees the first 32 bits of a 64 bit value from one write, and the second 32 bits from another write."

# Parallelizing Floyd's Algorithm

Which loops are parallelizable?

$$d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$$

- ▶ **for** (**int**  $r = 0$ ;  $r < n$ ;  $++r$ ) { ... }
- ▶ This loop can be parallelized.



Row slicing

```
for  $i = 0$  to  $n - 1$ 
  pfor  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```


- ▶ Any load-balancing issues?

# Parallelizing Floyd's Algorithm

Which loops are parallelizable?

$$d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$$

- ▶ **for** (**int**  $c = 0$ ;  $c < n$ ;  $++c$ ) { ... }
- ▶ This loop can be parallelized. (Same analysis as before.)



Column slicing

```
for  $i = 0$  to  $n - 1$ 
  for  $r = 0$  to  $n - 1$ 
    pfor  $c = 0$  to  $n - 1$ 
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

- ▶ Any load-balancing issues?

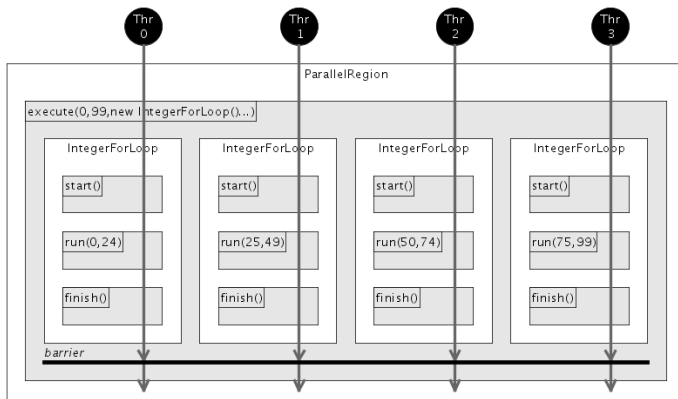
# FloydSmpRow.java

code/FloydSmpRow.java



# FloydSmpRow.java


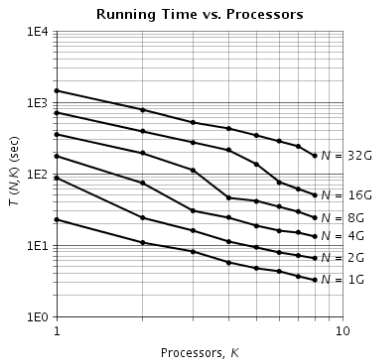
code/FloydSmpRow.java





# FloydSmpAltRow.java

code/FloydSmpAltRow.java


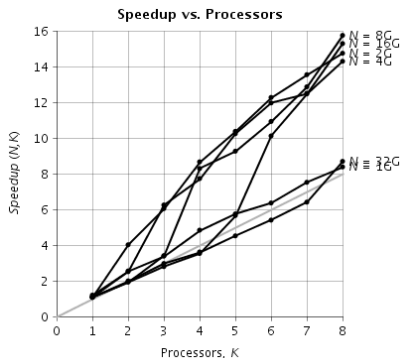
# FloydSmpRow Running Time and EDSF



# FloydSmpRow Speedup and Efficiency





# FloydSmpRow Speedup and Efficiency



The “unfair” JIT-compiler effect:

- ▶ FloydSmpRow has a “hot” `IntegerForLoop.run()` method.
- ▶ FloydSeq has a “naked” `for`-loop.

# FloydSmpRow Speedup and Efficiency



The “unfair” JIT-compiler effect:

- ▶ FloydSmpRow has a “hot” `IntegerForLoop.run()` method.
- ▶ FloydSeq has a “naked” `for`-loop.

More than “just” the JIT-compiler effect.

## FloydSmpRow Results

An abrupt jump in efficiencies as  $K$  increases.

- Larger  $N$  requires greater  $K$  before jump.

```
for  $i = 0$  to  $n - 1$ 
  pfor  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

The distance matrix requires  $\approx 8n^2$  bytes.

Each thread accesses only  $\approx 8(\frac{n^2}{K} + n)$  bytes per  $i$  iteration.

Futhermore, each thread accesses the *same*  $\approx 8\frac{n^2}{K}$  bytes each  $i$  iteration.

## FloydSmpRow Results

An abrupt jump in efficiencies as  $K$  increases.

- Larger  $N$  requires greater  $K$  before jump.

```
for  $i = 0$  to  $n - 1$ 
  pfor  $r = 0$  to  $n - 1$ 
    for  $c = 0$  to  $n - 1$ 
       $d_{rc} \leftarrow \min(d_{rc}, d_{ri} + d_{ic})$ 
```

The distance matrix requires  $\approx 8n^2$  bytes.

Each thread accesses only  $\approx 8(\frac{n^2}{K} + n)$  bytes per  $i$  iteration.

Futhermore, each thread accesses the *same*  $\approx 8\frac{n^2}{K}$  bytes each  $i$  iteration.



When threads' data all fits in their processors' cache,  
parallel program performs *much* better than sequential,  
which suffers from continual cache churning.





# FloydSmpCol.java

code/FloydSmpCol.java

# FloydSmpColSmp Running Time and EDSF



# FloydSmpColSmp Speedup and Efficiency



## FloydSmpCol Results

Classic Amdahl's Law behavior.

- ▶ Speedups approaching a limit
- ▶ Efficiencies continually decreasing as  $K$  increases
- ▶ Constant sequential fraction

In fact, a very large sequential fraction.

Why does FloydSmpCol have a larger sequential fraction?

## FloydSmpCol Results

Classic Amdahl's Law behavior.

- ▶ Speedups approaching a limit
- ▶ Efficiencies continually decreasing as  $K$  increases
- ▶ Constant sequential fraction

In fact, a very large sequential fraction.

Why does FloydSmpCol have a larger sequential fraction?

FloydSmpCol requires  $n^2$  barrier waits,  
but FloydSmpRow requires only  $n$  barrier waits.

Parallelizing the innermost loop typical yields poor performance.

## FloydSimpleRev{Seq,SmpCol}Method.java

What if we reverse the loops?

```
code/FloydSimpleSeqMethod.java  
code/FloydSimpleRevSeqMethod.java  
code/FloydSimpleRevSmpColMethod.java
```

Now, only  $n$  barrier waits.

As before, each thread accesses only  $\approx 8(\frac{n^2}{K} + n)$  bytes per  $i$  iteration and each thread accesses the *same*  $\approx 8\frac{n^2}{K}$  bytes each  $i$  iteration.



Same results?

# Cellular Automata

A cellular automaton (CA) is a simple abstract computing device that is capable of generating many kinds of interesting behavior.

The state of the system consists of a regular grid of cells (each of which has a value). At each time step, each cell simultaneously changes its value based on the values of a neighborhood of cells (according to some fixed rule).

A one-dimensional cellular automaton (1-D CA) uses an array of cells and the neighborhood of a cell consists of the cell to the left, the cell itself, and the cell to the right (using wraparound boundaries).



# Elementary Cellular Automata

An elementary cellular automaton (ECA)  
is a one-dimensional discrete cellular automaton (1-D DCA)  
where each cell has a value that is either 0 or 1.

The rule for updating cells can be represented as an unsigned 8-bit integer

- interpret the left, center, and right cells as an unsigned 3-bit integer called  $n$   
and the new state of the center cell is the  $n^{\text{th}}$  bit of the 8-bit rule.

For instance, rule **30** = **00011110**<sub>2</sub> corresponds to the following:


configuration of cells	111	110	101	100	011	010	001	000
new state of center cell	0	0	0	1	1	1	1	0



# Elementary Cellular Automata

For instance, rule **30** =  $00011110_2$  corresponds to the following:

configuration of cells	111	110	101	100	011	010	001	000
new state of center cell	0	0	0	1	1	1	1	0



500 cells, 250 steps, rule **30** =  $00011110_2$

## ElementaryCA{Seq,SmpAlt}.java

```
boolean[][] cells = new boolean[S+1][N];
cells[0][c/2] = true;
for (int s = 1; s <= S; s++)
    for (int c = 0; c < N; c++)
        cells[s][c] = applyRule(cells[s-1][c-1],
                                cells[s-1][c],
                                cells[s-1][c+1]);

int total = 0;
for (int c = 0; c < N; c++)
    if (cells[S][c]) total++;
```

Requires  $O(SN)$  memory to hold data;  
problematic when scaling up the problem.

## ElementaryCA{Seq,SmpAlt}.java

Only need previous row to calculate next row.

```
boolean[] next = new boolean[N];
boolean[] cells = new boolean[N];
cells[c/2] = true;
for (int s = 1; s <= S; s++) {
    for (int c = 0; c < N; c++)
        next[c] = applyRule(cells[c-1], cells[c], cells[c+1]);
    boolean[] temp = cells;
    cells = next;
    next = temp;
}
for (int c = 0; c < N; c++)
    if (cells[c]) total++;
```

Requires only  $O(N)$  memory to hold data.

# ElementaryCA{Seq,SmpAlt}.java

code/ElementaryCASeq.java  
code/ElementaryCASmpAlt.java

## Parallelizing ElementaryCA:

- ▶ What portions of the Elementary CA program are not parallelizable?
- ▶ What portions of the Elementary CA program are parallelizable?
  - ▶ What parallelization patterns?
- ▶ Any synchronization issues?
- ▶ Any load-balancing issues?
- ▶ Any cache interference?
- ▶ Any beneficial cache effects?

## Continuous Cellular Automata

Another CA is a one-dimensional continuous cellular automaton (1-D CCA) where each cell has a value that is a rational number in the range **0** to **1**.

The rule for updating cells uses two rational constants **A** and **B**:

$$x_i^{\text{new}} = \text{frac} \left( \frac{x_{i-1} + x_i + x_{i+1}}{3} \cdot A + B \right)$$

# Continuous Cellular Automata

$$x_i^{\text{new}} = \text{frac} \left( \frac{X_{i-1} + X_i + X_{i+1}}{3} \cdot A + B \right); A = 1, B = 11/12$$

s	x <sub>0</sub>	x <sub>1</sub>	x <sub>2</sub>	x <sub>3</sub>	x <sub>4</sub>	x <sub>5</sub>	x <sub>6</sub>	x <sub>7</sub>	x <sub>8</sub>	x <sub>9</sub>
0	0	0	0	0	0	1	0	0	0	0

# Continuous Cellular Automata

$$x_i^{\text{new}} = \text{frac} \left( \frac{X_{i-1} + X_i + X_{i+1}}{3} \cdot A + B \right); A = 1, B = 11/12$$

s	x <sub>0</sub>	x <sub>1</sub>	x <sub>2</sub>	x <sub>3</sub>	x <sub>4</sub>	x <sub>5</sub>	x <sub>6</sub>	x <sub>7</sub>	x <sub>8</sub>	x <sub>9</sub>
0	0	0	0	0	0	1	0	0	0	0
1	11/12	11/12	11/12	11/12	1/4	1/4	1/4	11/12	11/12	11/12

# Continuous Cellular Automata

$$x_i^{\text{new}} = \text{frac} \left( \frac{X_{i-1} + X_i + X_{i+1}}{3} \cdot A + B \right); A = 1, B = 11/12$$


s	x <sub>0</sub>	x <sub>1</sub>	x <sub>2</sub>	x <sub>3</sub>	x <sub>4</sub>	x <sub>5</sub>	x <sub>6</sub>	x <sub>7</sub>	x <sub>8</sub>	x <sub>9</sub>
0	0	0	0	0	0	1	0	0	0	0
1	11/12	11/12	11/12	11/12	1/4	1/4	1/4	11/12	11/12	11/12
2	5/6	5/6	5/6	11/18	7/18	1/6	7/18	11/18	5/6	5/6
3	3/4	3/4	73/108	19/36	11/36	25/108	11/36	19/36	73/108	3/4
4	2/3	52/81	46/81	34/81	22/81	16/81	22/81	34/81	46/81	52/81
5	$\frac{551}{972}$	$\frac{527}{972}$	$\frac{149}{324}$	$\frac{109}{324}$	$\frac{23}{108}$	$\frac{53}{324}$	$\frac{23}{108}$	$\frac{109}{324}$	$\frac{149}{324}$	$\frac{527}{972}$




# Continuous Cellular Automata

$$x_i^{\text{new}} = \text{frac} \left( \frac{X_{i-1} + X_i + X_{i+1}}{3} \cdot A + B \right); A = 1, B = 11/12$$


s	x <sub>0</sub>	x <sub>1</sub>	x <sub>2</sub>	x <sub>3</sub>	x <sub>4</sub>	x <sub>5</sub>	x <sub>6</sub>	x <sub>7</sub>	x <sub>8</sub>	x <sub>9</sub>
0	0	0	0	0	0	1	0	0	0	0
1	11/12	11/12	11/12	11/12	1/4	1/4	1/4	11/12	11/12	11/12
2	5/6	5/6	5/6	11/18	7/18	1/6	7/18	11/18	5/6	5/6
3	3/4	3/4	73/108	19/36	11/36	25/108	11/36	19/36	73/108	3/4
4	2/3	52/81	46/81	34/81	22/81	16/81	22/81	34/81	46/81	52/81
5	551/972	527/972	149/324	109/324	23/108	53/324	23/108	109/324	149/324	527/972



# Continuous Cellular Automata



400 cells, 200 steps,  $A = 1$ ,  $B = 11/12$



400 cells, 200 steps,  $A = 13/12$ ,  $B = 11/12$

# Continuous Cellular Automata

Use rational arithmetic, not floating-point arithmetic.

- ▶ Floating-point arithmetic does not have sufficient precision; rounding errors would quickly accumulate and lead to incorrect results.
- ▶ `edu.rit.numeric.BigRational`
  - ▶ Represent numerator and denominator with arbitrary precision integers (`java.math.BigInteger`).
  - ▶ Convert to **float** or **double** via arbitrary precision decimals (`java.math.BigDecimal`).
    - ▶ (Necessary loss of precision when converting to 8-bit grayscale value.)

# Continuous Cellular Automata

Would like to produce *images*, not just a *reduction*.

Back to requiring  $O(SN)$  memory to store image?

# Continuous Cellular Automata

Would like to produce *images*, not just a *reduction*.

Back to requiring  $O(SN)$  memory to store image?

```
// Write all rows and columns of the image to the output stream.
```

```
void PJGImage.Writer.write();
```

```
// Write the given row slice of the image to the output stream.
```

```
void PJGImage.Writer.writeRowSlice(Range theRowRange);
```

```
// Write the given column slice of the image to the output stream.
```

```
void PJGImage.Writer.writeColSlice(Range theColRange);
```

```
// Write the given patch of the image to the output stream.
```

```
void PJGImage.Writer.writePatch(Range theRowRange, Range theColRange);
```

## CCASeq.java

code/CCASeq.java

Why doesn't `static byte[][] pixelmatrix` lead to  $O(SN)$  memory?

code/CCASeq.java

Why doesn't `static byte[][] pixelmatrix` lead to  $O(SN)$  memory?

Parallelizing CCA:

- ▶ What portions of the Continuous CA program are not parallelizable?
- ▶ What portions of the Continuous CA program are parallelizable?
  - ▶ What parallelization patterns?
- ▶ Any synchronization issues?
- ▶ Any load-balancing issues?
- ▶ Any cache interference?
- ▶ Any beneficial cache effects?

CCASmpAlt.java


code/CCASmpAlt.java



## Barrier Actions

```
new ParallelTeam().execute(new ParallelRegion() {  
    public void run() {  
        ...  
        execute(0, 99,  
            new IntegerForLoop() {  
                public void run(int first, int last) {  
                    for (int i = first; i <= last; i++) {  
                        ... // Loop body  
                    }  
                }  
            },  
            new BarrierAction() {  
                public void run() {  
                    ... // Code to be executed in a single thread  
                }  
            }  
        });  
        ...  
    }  
});
```

# Barrier Actions



## Other Barrier Actions

```
// Instead of a barrier action object,  
// use a constant BarrierAction:
```

```
// Each thread waits at the barrier.  
execute(0, 99, new IntegerForLoop { ... },  
        BarrierAction.WAIT);
```

```
// Each thread does not wait at the barrier.  
execute(0, 99, new IntegerForLoop { ... },  
        BarrierAction.NO_WAIT);
```

In what situations would `NO_WAIT` be useful?

- ▶ Remember: correctness trumps performance

CCASmp.java



code/CCASmp.java

CCASmp.java



code/CCASmp.java

What is the advantage of `CCASmp.java` over `CCASmpAlt.java`?

# CCASmp Running Time and EDSF




# CCASmp Speedup and Efficiency



Classic Amdahl's Law behavior,  
but with a rather large sequential fraction.

What portions of the Continuous CA program  
are causing the large sequential fraction?

# Beating Amdahl's Law





# Beating Amdahl's Law with Overlapping

What portions of the Continuous CA program are causing the large sequential fraction?

The barrier action that

- ▶ computes grayscale value of cell state
- ▶ writes pixel row to image file

## Beating Amdahl's Law with Overlapping

What portions of the Continuous CA program are causing the large sequential fraction?

The barrier action that

- ▶ computes grayscale value of cell state
- ▶ writes pixel row to image file

(Any missed parallelism within these actions?)

## Beating Amdahl's Law with Overlapping

What portions of the Continuous CA program are causing the large sequential fraction?

The barrier action that

- ▶ computes grayscale value of cell state
- ▶ writes pixel row to image file

(Any missed parallelism within these actions?)

Any missed parallelism between these actions and next-state computation?

## Beating Amdahl's Law with Overlapping

What portions of the Continuous CA program are causing the large sequential fraction?

The barrier action that

- ▶ computes grayscale value of cell state
- ▶ writes pixel row to image file

(Any missed parallelism within these actions?)

Any missed parallelism between these actions and next-state computation?

- ▶ Could write the current state to image file while computing the next state.

## Beating Amdahl's Law with Overlapping

What portions of the Continuous CA program are causing the large sequential fraction?

The barrier action that

- ▶ computes grayscale value of cell state
- ▶ writes pixel row to image file


(Any missed parallelism within these actions?)

Any missed parallelism between these actions and next-state computation?


- ▶ Could write the current state to image file while computing the next state.

Overlapping: run the I/O thread in parallel with the computation threads.

# Beating Amdahl's Law with Overlapping



# Beating Amdahl's Law with Overlapping



Are we *really* beating Amdahl's Law?

# Overlapping: Speedup and Efficiency

Running Time

$$\blacktriangleright T(N, K) = \max( F \cdot T(N, 1) , \frac{1}{K} \cdot (1 - F) \cdot T(N, 1) )$$

Speedup



# Overlapping: Speedup and Efficiency

Running Time

$$\blacktriangleright T(N, K) = \max( F \cdot T(N, 1) , \frac{1}{K} \cdot (1 - F) \cdot T(N, 1) )$$

Speedup

$$\blacktriangleright \text{Speedup}(N, K) = \frac{T(N, 1)}{T(N, K)} = \min( \frac{1}{F} , \frac{K}{1-F} )$$

Efficiency

# Overlapping: Speedup and Efficiency

Running Time

$$\blacktriangleright T(N, K) = \max( F \cdot T(N, 1) , \frac{1}{K} \cdot (1 - F) \cdot T(N, 1) )$$

Speedup

$$\blacktriangleright \text{Speedup}(N, K) = \frac{T(N, 1)}{T(N, K)} = \min( \frac{1}{F} , \frac{K}{1-F} )$$

Efficiency

$$\blacktriangleright \text{Eff}(N, K) = \frac{\text{Speedup}(N, K)}{K} = \min( \frac{1}{K \cdot F} , \frac{1}{1-F} )$$

# Overlapping: Speedup and Efficiency

Running Time

$$\blacktriangleright T(N, K) = \max( F \cdot T(N, 1) , \frac{1}{K} \cdot (1 - F) \cdot T(N, 1) )$$

Speedup

$$\blacktriangleright \text{Speedup}(N, K) = \frac{T(N, 1)}{T(N, K)} = \min( \frac{1}{F} , \frac{K}{1-F} )$$


Efficiency

$$\blacktriangleright \text{Eff}(N, K) = \frac{\text{Speedup}(N, K)}{K} = \min( \frac{1}{K \cdot F} , \frac{1}{1-F} )$$

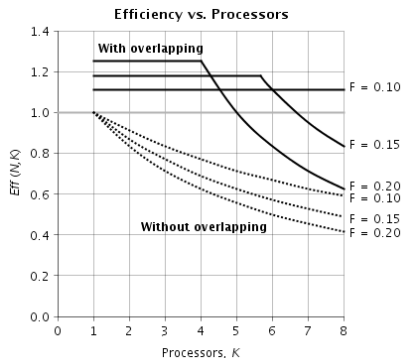
What happens to *Speedup* and *Eff* for small  $K$ ?

What happens to *Speedup* and *Eff* as  $K \rightarrow \infty$ ?

# Overlapping: Speedup and Efficiency



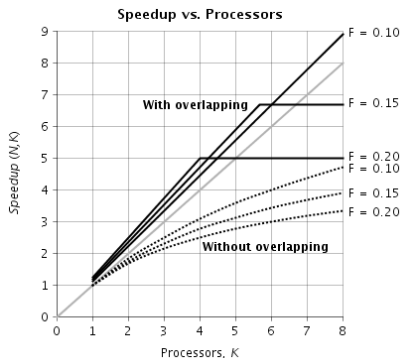
$$\min\left(\frac{1}{F}, \frac{K}{1-F}\right)$$




$$\min\left(\frac{1}{F \cdot K}, \frac{1}{1-F}\right)$$

In the limit, as  $K \rightarrow \infty$ ?

# Overlapping: Speedup and Efficiency



$$\min\left(\frac{1}{F}, \frac{K}{1-F}\right)$$



$$\min\left(\frac{1}{F \cdot K}, \frac{1}{1-F}\right)$$

In the limit, as  $K \rightarrow \infty$ ?

I claim: this superlinear speedup is “cheating”. Why?

# Overlapping: Speedup and Efficiency

Running Time

$$\blacktriangleright T(N, K) = \max( F \cdot T(N, 1), \frac{1}{K-1} \cdot (1 - F) \cdot T(N, 1) )$$

Speedup

$$\blacktriangleright \text{Speedup}(N, K) = \frac{T(N, 1)}{T(N, K)} = \min( \frac{1}{F}, \frac{K-1}{1-F} )$$

Efficiency

$$\blacktriangleright \text{Eff}(N, K) = \frac{\text{Speedup}(N, K)}{K} = \min( \frac{1}{K \cdot F}, \frac{1}{1-F} - \frac{1}{K \cdot (1-F)} )$$

What happens to *Speedup* and *Eff* for small  $K$ ?

What happens to *Speedup* and *Eff* as  $K \rightarrow \infty$ ?

# Overlapping: Speedup and Efficiency

Running Time

$$\blacktriangleright T(N, K) = \max( F \cdot T(N, 1) , \frac{1}{K-1} \cdot (1 - F) \cdot T(N, 1) )$$

Speedup

$$\blacktriangleright \text{Speedup}(N, K) = \frac{T(N, 1)}{T(N, K)} = \min( \frac{1}{F} , \frac{K-1}{1-F} )$$

Efficiency

$$\blacktriangleright \text{Eff}(N, K) = \frac{\text{Speedup}(N, K)}{K} = \min( \frac{1}{K \cdot F} , \frac{1}{1-F} - \frac{1}{K \cdot (1-F)} )$$

What happens to *Speedup* and *Eff* for small  $K$ ?

What happens to *Speedup* and *Eff* as  $K \rightarrow \infty$ ?

No superlinear speedup.

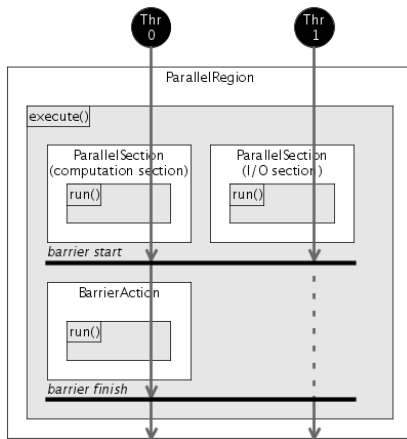
## Parallel Sections

Run a parallel team of threads,  
where each thread may execute different code.

```
new ParallelTeam(2).execute(new ParallelRegion() {  
    public void run() {  
        execute (  
            new ParallelSection() {  
                public void run() {  
                    // Code for computation  
                }  
            },  
            new ParallelSection() {  
                public void run() {  
                    // Code for I/O  
                }  
            },  
            new BarrierAction() {  
                public void run() {  
                    // Code for single-threaded barrier action  
                }  
            }  
        ));  
    }  
});
```



# Parallel Sections

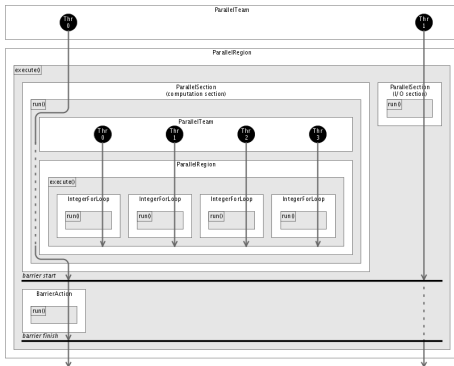


## Nested Parallel Regions

In the 1-D CCA program, the computation task is a parallel problem. The computation section contains another (nested) parallel region.

# Nested Parallel Regions



In the 1-D CCA program, the computation task is a parallel problem. The computation section contains another (nested) parallel region.





CCASmp2.java

code/CCASmp2.java

# CCASmp2 Running Time and EDSF





# CCASmp2 Running Time and EDSF





Why the spike at  $K = 8$ ?

# CCASmp2 Speedup and Efficiency





# CCASmp2 Speedup and Efficiency



Why the dip at  $K = 8$ ?



# CCASmp vs. CCASmp2 Running Time and EDSF



# CCASmp vs. CCASmp2 Speedup and Efficiency

