

**4003-531/4005-735
Term 20102**

Parallel Computing I

**Assignment 3
January 19, 2011**

**Assignment 3
Due: February 2, 2011**

Before starting, be sure to review the instructions for running Parallel Java on the RIT CS Parallel Computers (<http://www.cs.rit.edu/~ark/runningpj.shtml>). Remember, to execute programs on the `thugN` backend nodes, submit it to the job queue running on `paranoia`. **Directly logging into the `thugN` backend nodes is not allowed.**

1. (30pts)

Building Parallel Programs, Part III Exercises:

- (a) Exercise 21 (p. 574)
- (b) Exercise 22 (p. 574)
- (c) Exercise 23 (p. 574)
- (d) Exercise 24 (p. 574)
- (e) Exercise 25 (p. 574)
- (f) Exercise 26 (p. 574)

Submission Submit a plain text file named `hw3-1.txt` or a PDF file named `hw3-1.pdf`. The `hw3-1` file should contain the solutions to the exercises.

2. (30pts)

Review the elementary cellular automata problem from Assignment 1. You may download the reference solution files `ElementaryCAsEq.java` and `ElementaryCAsMp.java` from the course website.

Recall the following:

- The `ElementaryCAsEq` program requires $O(\text{gridSize} \cdot \text{numSteps} + \text{gridSize})$ computation.
- The `ElementaryCAsEq` and `ElementaryCAsMp` programs allocated exactly two arrays of size `gridSize`.
- The `ElementaryCAsMp` program used a parallel `for`-loop to parallelize the update of the cellular automaton state and a parallel `for`-loop to parallelize the summation of the number of 1 cells in the final configuration.
- A fixed schedule sufficed for the `ElementaryCAsMp` program, because each individual cell update requires the same computation.

Write Java programs called `ElementaryCAsEq`¹ (a sequential program) and `ElementaryCAClu` (a parallel program) that executes elementary cellular automata. Both programs must take the following command-line arguments:

- *rule*: the rule to execute (an integer)
- *gridSize*: the size of the grid (an integer)
- *numSteps*: the number of steps to execute (an integer)

The initial state of the cellular automaton is a grid of `gridSize` cells, where exactly one cell is 1 and all other cells are 0. When executing the cellular automaton, assume that the grid “wraps around”; that is, the cell to the left of the first cell of the grid (i.e., at index 0) is the last cell of the grid (i.e., at index `gridSize - 1`) and the cell to the right of the last cell of the grid (i.e., at index `gridSize - 1`) is the first cell of the grid (i.e., at index 0). After executing the elementary cellular automaton for `numSteps` steps, the program should print out the number of 1 cells in the final configuration and the running time of the program; in the cluster parallel program, only rank 0 should print out the total number of 1 cells in the final configuration, but each rank should print out the running time of the program. For instance, here are a sample executions of `ElementaryCAsEq` and `ElementaryCAClu`:

```
[mtf@paranoia code]$ java ElementaryCAsEq 30 200000 2000
Job 49, thug20
2021
Running time: 33414 msec
[mtf@paranoia code]$ java -Dpj.np=4 ElementaryCAClu 30 200000 2000
Job 50, thug21, thug22, thug23, thug24
Running time: 11017 msec; rank(3)
Running time: 11070 msec; rank(2)
Running time: 11085 msec; rank(1)
2021
Running time: 11081 msec; rank(0)
```

¹You may use the provided reference solution, but a sequential program should begin timing *before* calling `Comm.init` when it is being compared to a cluster parallel program.

One advantage of cluster parallel computers over SMP parallel computers is that they may scale to much larger total memory; hence, with a cluster parallel computer, one may solve problems that require more total memory than can be accommodated in a single SMP parallel computer. One requirement of the `ElementaryCAClu.java` program is that it should not allocate (significantly) more than two arrays of size $gridSize/K$, where K is the number of processes in the cluster parallel program.

Considering `ElementaryCAClu`, give an expression for the amount of data (measured in bits) sent during the execution of the program. This formula should be a function of $gridSize$, $numSteps$, and K . Be sure to note if there are any special cases. You may ignore the communication-layer overhead.²

Measure the running time T for `ElementaryCASeq` with command-line arguments `90 gridSize numSteps` on the `thugN` backend nodes and `ElementaryCAClu` with the same command-line arguments on the `thugN` backend nodes using 1, 2, 4, and 8 processors. Calculate *Speedup*, *Eff*, and *EDSF*. When measuring the running time T for `ElementaryCAClu`, be sure to record the running time of the rank which has the longest running time. Submit your results in a table organized as follows:

<i>gridSize</i>	<i>numSteps</i>	<i>K</i>	<i>T</i> (msec)	<i>Speedup</i>	<i>Eff</i>	<i>EDSF</i>
20000000	200	seq		xxx	xxx	xxx
20000000	200	1				xxx
20000000	200	2				
20000000	200	4				
20000000	200	8				
20000000	200	16				
2000000	2000	seq		xxx	xxx	xxx
2000000	2000	1				xxx
2000000	2000	2				
2000000	2000	4				
2000000	2000	8				
2000000	2000	16				
20000	200000	seq		xxx	xxx	xxx
20000	200000	1				xxx
20000	200000	2				
20000	200000	4				
20000	200000	8				
20000	200000	16				

Explain your results.

Submission Submit `ElementaryCASeq.java`, `ElementaryCAClu.java`, and a plain text file named `hw3-2.txt` or a PDF file named `hw3-2.pdf`. The `hw3-2` file should contain the formula for amount of data sent during the execution of the cluster parallel program, the tabulated results, and an explanation.

²That is, if your program performs a `world.send(toRank, intBuf)`, where `intBuf` is a `IntegerItemBuf`, then count that as 32 bits of data sent (one `int`), and not as 32 bits plus Parallel Java header bits plus TCP header bits plus IP header bits plus Note that Parallel Java sends a `boolean` as one (8 bits) byte.

Submission

Submit a single ZIP file named `hw3.zip` to the Homework 3 Dropbox on MyCourses by the due date. The `hw3.zip` file should contain:

- `hw3-1.txt` or `hw3-1.pdf`
- `ElementaryCASeq.java`
- `ElementaryCAClu.java`
- `hw3-2.txt` or `hw3-2.pdf`

The `hw3.zip` file should contain no additional files.

Document History

January 19, 2011

Original version